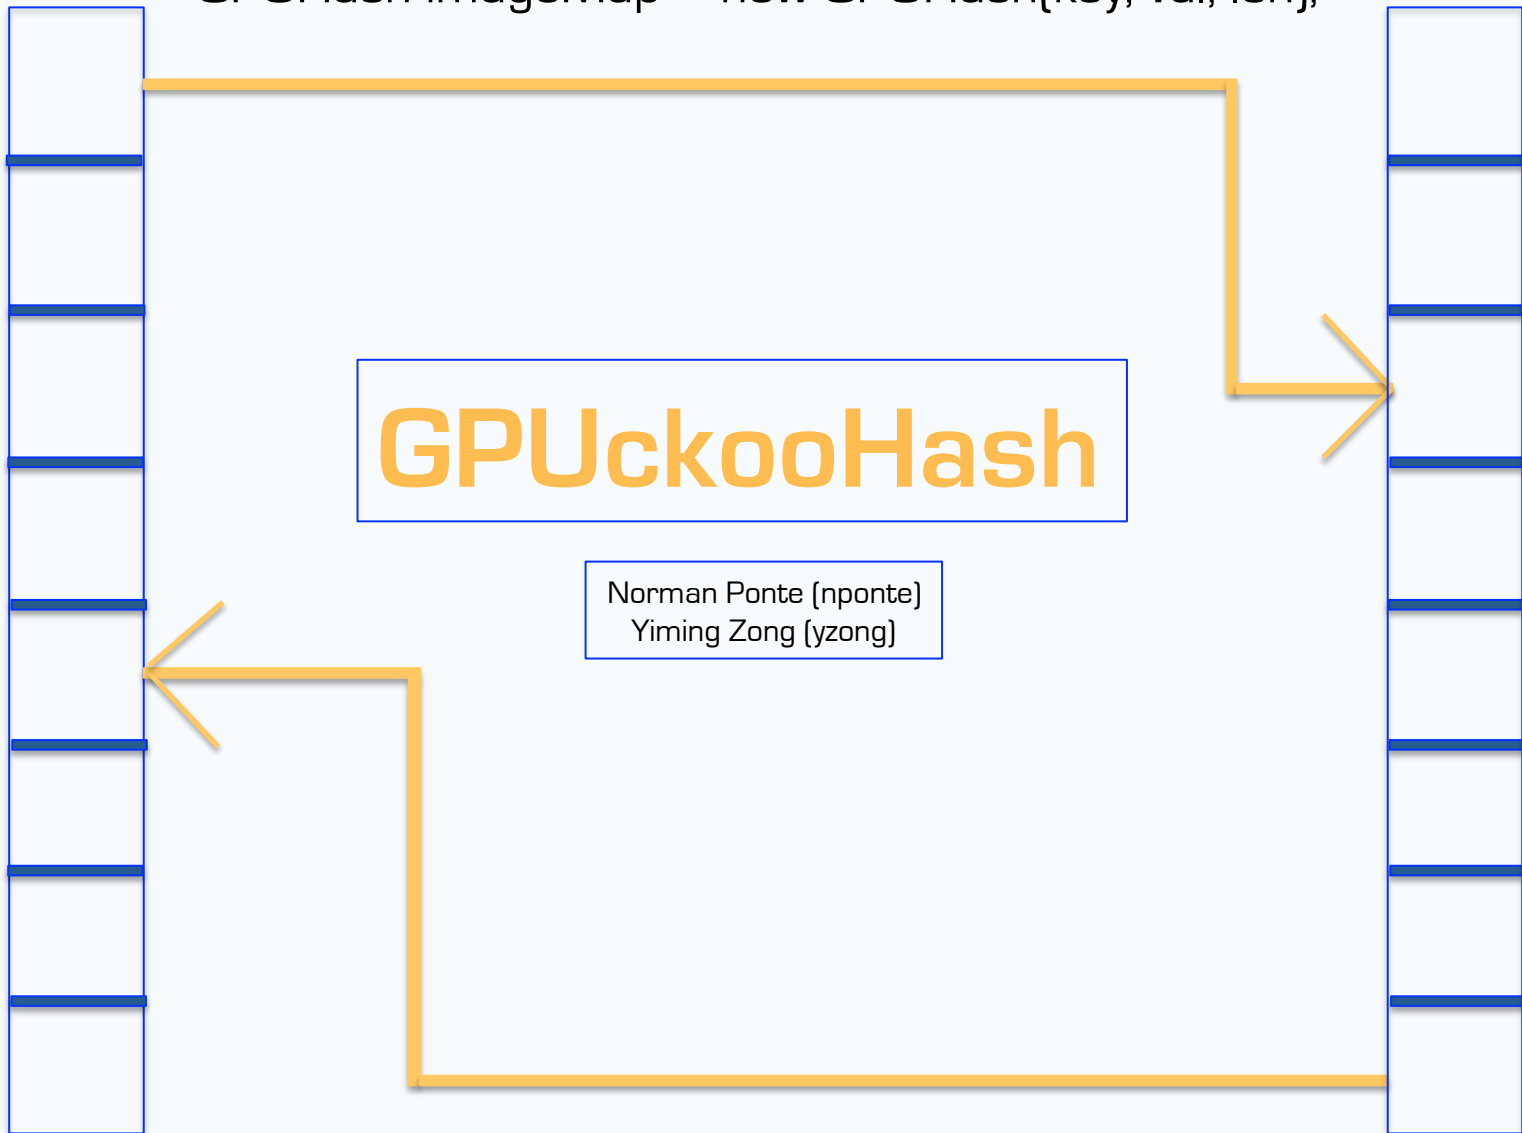


```
#import "GPUHash.h"
```

```
GPUHash imageMap = new GPUHash(key, val, len);
```



M

PH?

CC

H

D

B

BS

MvS

B.I

B.II

B.III

* (Motivation)


- Hash large images to short numbers
- Similar characteristic hash similarly
- Hamming Distance – difference in bits between two hashes
- Spatial Hashing




4,000 / sec



58,000,000 / day



376 TB



2014 - 880,000,000,000

M
PH?
CC
H
D
B
BS
MvS
B.I
B.II
B.III

Parallel Hashing?

- Synchronization
- Variable workloads
- High memory bandwidth
- Large data structures

sparsehash / sparsehash

<> Code
🔔 Issues 8
🔗 Pull requests 0
📶 Pulse
📊 Graphs

Automatically exported from code.google.com/p/sparsehash

RESOURCE USAGE


* sparse_hash_map has memory overhead of about 2 bits per hash-map entry.

* dense_hash_map has a factor of 2-3 memory overhead: if your hashtable data takes X bytes, dense_hash_map will use 3X-4X memory total.

hash(key1, val1) – Bucket 1 -

hash(key2, val2) – Bucket 2 -

hash(key3, val3) – Bucket 1 -



Xeon Phi L1 Cache :: 64 bytes

M

PH?

CC

H

D

B

BS

MvS

B.I

B.II

B.III

Optimizing for Our Common Case

3 Metrics : Speed, Memory Footprint, Construction Time


- Care about mostly lookup
- Want rapid construction and deletion
- User defined Space vs. Speed tradeoff
- Batch Processing
- CUDA thrust library

```

/** @brief Look into hash table and return values for
 *      given keys.
 *
 *      * Probe the hash table and return value for input    a
 *      * array key. 0 means failure.
 *
 *      * @param keys – keys array to lookup
 *      * @param val_ret – array to fill with values for
 *      *                  corresponding keys
 *      * @param len – length of keys
 *      * @return Void.
 *      */
void GPUHash::lookupTable(int* key, int* val_ret, const int len)

```

Thrust Performance vs. Intel TBB



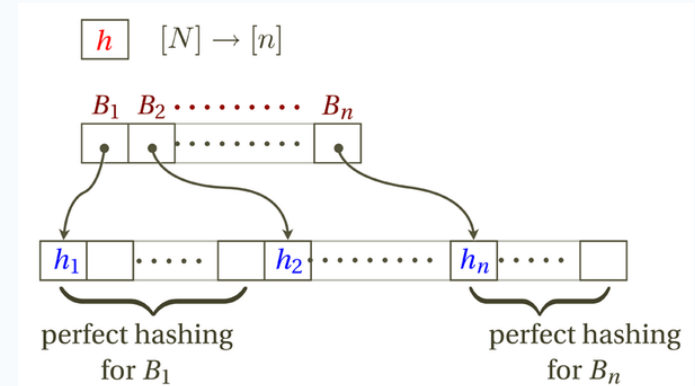
Performance may vary based on OS version and motherboard configuration
 • Thrust v1.7.1 on K40m, ECC ON, input and output data on device
 • TBB 4.2 on Intel hybrid 12-core E5-2697 v2 @ 2.70GHz

M
PH?
CC
H
D
B
BS
MvS
B.I
B.II
B.III


FKS Hashing + Cuckoo Hashing

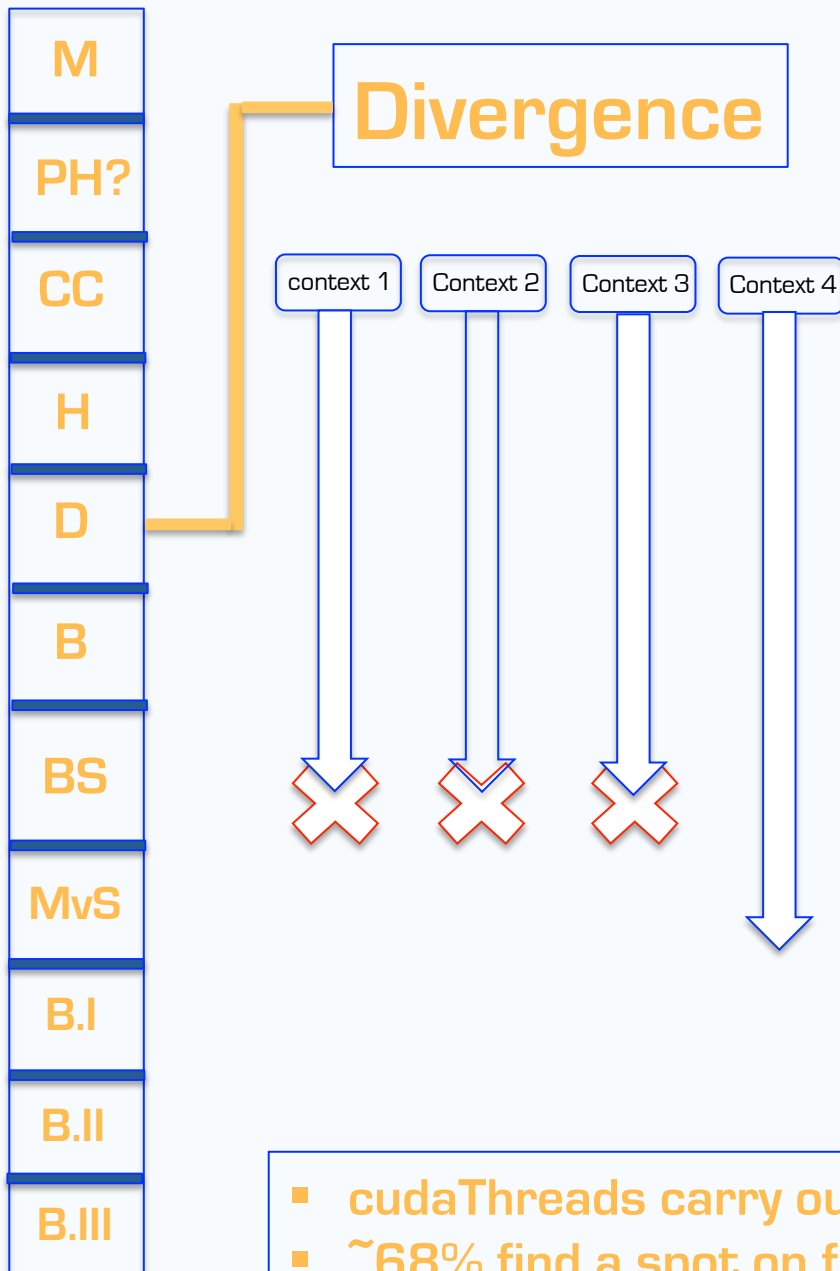
- Lock Free
- $O(1)$ - Lookup time
- 3 - Seed Cuckoo Hashing
- Hash to FKS Bucket
- Hash to 3 Bucket inside FKS Bucket using cuckoo hash functions
- Random Seed

FKS Hashing



Cuckoo Hashing






```
// Loop for all cuckoo hashing functions
for (int i = 0; i < NUM_FUNC; i++) {
    // Use one of the hash functions
    unsigned int hash;
    switch (i) {
        case 0: hash = hashKey(key, b->get_hash1());
                break;
        case 1: hash = hashKey(key, b->get_hash2());
                break;
        case 2: hash = hashKey(key, b->get_hash3());
                break;
    }

    //Bucket does atomic exchange
    newkeyval = b->add_slot(key, val, i, hash);
    // If I got into the spot then exit routine
    if (newkeyval == 0) {
        return thrust::make_pair(0,0);
    }
    key = get_key(newkeyval);
    val = get_val(newkeyval);
}
```

- cudaThreads carry out 3 stage cuckoo hashing
- ~68% find a spot on first hash



Balancing

Decrease Bucket Factor Low

- Waste of Space
- Faster runtime
- Divergence
- Poor cache usage


Increase Bucket Factor High

- Space Efficient
- Slow runtime
- Could total Reseed
- Contention for atomic variable

Load Factor Best Loop to Converge

.3	3
.5	5
.8	14
.9	30

- Modify buckets based on number of elements
- Reseed FKS buckets if one bucket is full >1%
- Good Hash Function 4x
- Loop iterations depends on load factor

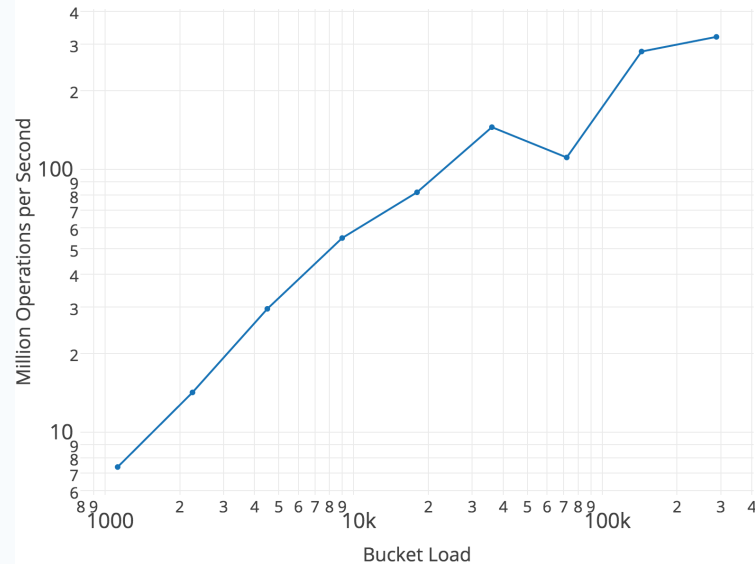


Bucket Size

BUCKET_LOAD = ?

- Bigger Bucket - can't fit in cache (512 KB)
- 80% time spent in sort
- Smaller bucket – pigeon hole principle
- 36 Byte Overhead per bucket
- Rebalancing cost is higher

Insertion / Deletion Throughput vs. Bucket Load



M

PH?

CC

H

D

B

BS

MvS

B.I

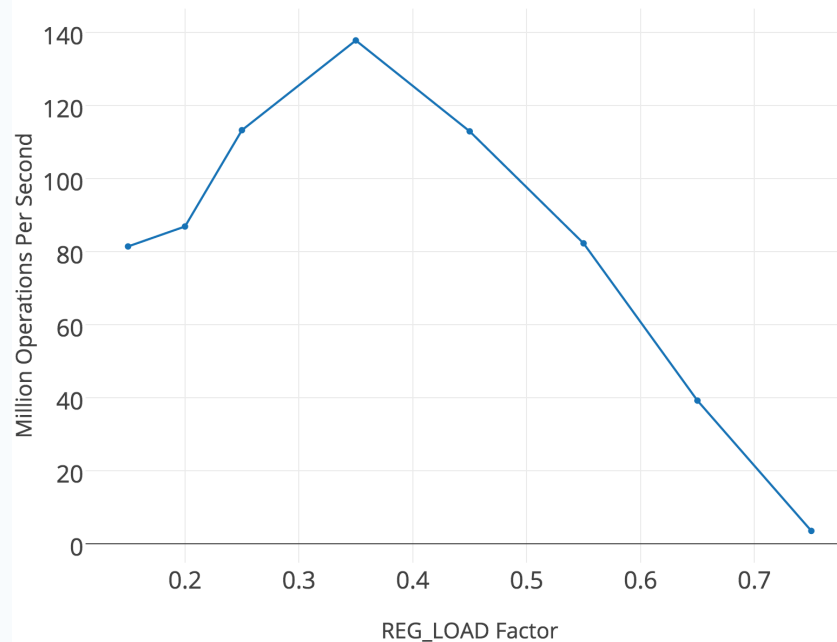
B.II

B.III

Memory and Speed

- Auto-rebalance for speed retention
- Perfect hashing N^2 space
- Permanent Memory Overhead O
- Want to remain near the peak

Insertion Throughput vs. REG_LOAD



M

PH?

CC

H

D

B

BS

MvS


B.I

B.II

B.III

Benchmark :: Lookup

Lookup Throughput vs. Data Size



■ Most important graph $O(1)$ – 3 Lookups

M

PH?

CC

H

D

B

BS

MvS

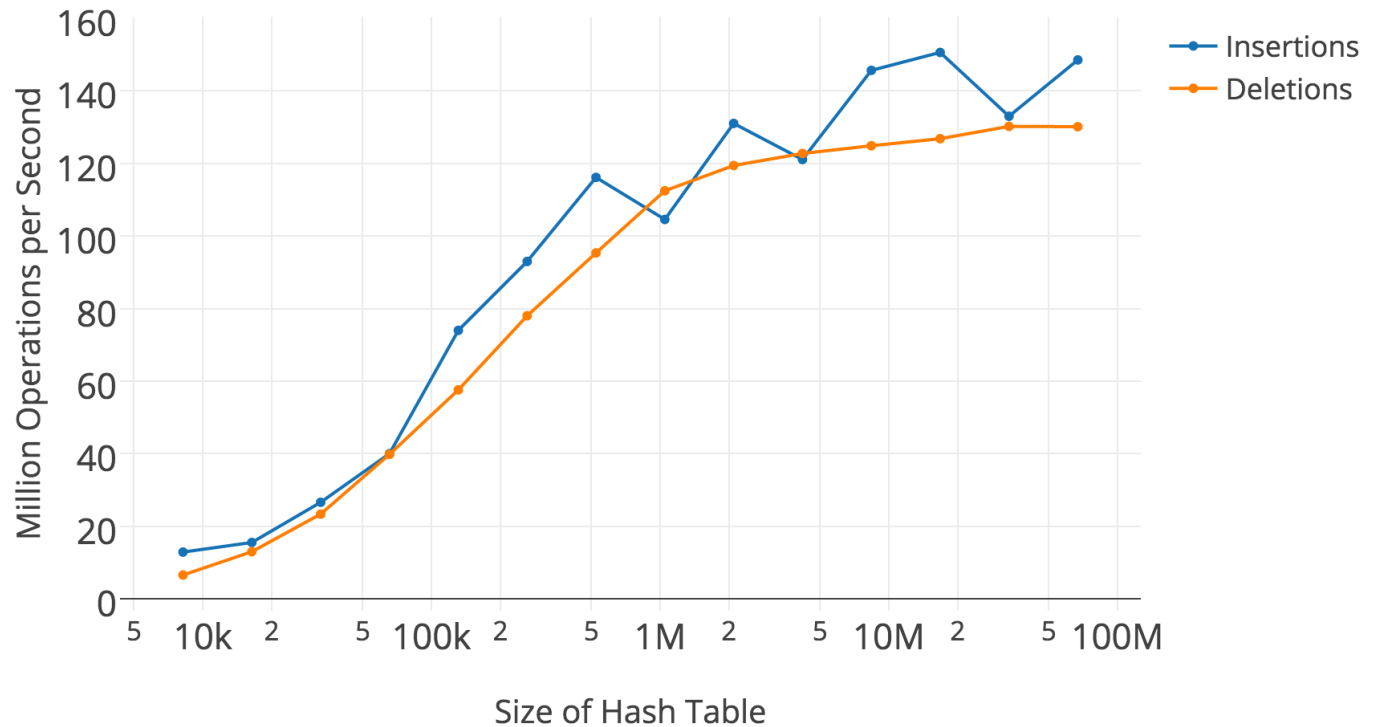
B.I

B.II

B.III

Benchmark :: insert + delete

Insertion / Deletion Throughput vs. Data Size



■ No rebalancing

M

PH?

CC

H

D

B

BS


MvS

B.I

B.II

B.III

Benchmark :: the field



- Considered to be one of the best lock free hash table
- Memory error

Q	u	e	s	?	t	i	o	n
---	---	---	---	---	---	---	---	---